18-447

Computer Architecture Lecture 10: Control Dependence Handling

Prof. Onur Mutlu
Carnegie Mellon University
Spring 2013, 2/8/2013

Reminder: Homework 2

- Homework 2 out
 - Due February 11 (Monday!)
 - LC-3b microcode
 - □ ISA concepts, ISA vs. microarchitecture, microcoded machines
- Remember: Homework 1 solutions were out

Reminder: Lab Assignment 2

- Lab Assignment 2
 - Due Feb 15 (next Friday!)
 - Single-cycle MIPS implementation in Verilog
 - All labs are individual assignments
 - No collaboration; please respect the honor code
 - Do not forget the extra credit portion!

Readings for Next Few Lectures

- P&H Chapter 4.9-4.11
- Smith and Sohi, "The Microarchitecture of Superscalar Processors," Proceedings of the IEEE, 1995
 - More advanced pipelining
 - Interrupt and exception handling
 - Out-of-order and superscalar execution concepts
- Recommended:
 - McFarling, "Combining Branch Predictors," DEC WRL Technical Report, 1993.

Also, ...

Kessler, "The Alpha 21264 Microprocessor," IEEE Micro 1999.

Monday: IEEE Tech Talk and CALCM Seminar

- Tuesday, February 12, 4:30-6:30pm, Hamerschlag 1107
- Dr. Richard E. Kessler, Cavium Fellow and Principal Architect
- Designing Efficient Processor Cores for Multicore Networking
- Abstract:
 - The design of CPUs has always required a balance of performance and efficiency in power, area, and complexity. The emergence of multicore SoCs armed with accelerators for packet processing has shifted this balance from solely single-thread performance to a combination of single-thread performance and efficient parallel processing. This shift requires a new style of core with short and deterministic pipelines, caches and memory systems optimized for low latency and high bandwidth, and an architecture that scales to 48-plus cores on a chip. This talk demonstrates how continuously emerging application demands shaped the fundamental principles behind OCTEON processor cores and supporting on-chip accelerators.

Monday: IEEE Tech Talk and CALCM Seminar

- Monday, February 11, 4:30-6:30pm, Hamerschlag 1107
- Dr. Richard E. Kessler, Cavium Fellow and Principal Architect
- Designing Efficient Processor Cores for Multicore Networking

Break-out Session:

Cavium is building a community of university and industry partners around the 32-core OCTEON II solution, with evaluation boards in use by students and professors at several universities globally. This break-out session for students will be conducted at the conclusion of the talk above to describe the evaluation board hardware, the Cavium SDK, and various semester-long student projects appropriate for upper level undergraduates or first year masters students. Other aspects of the OCTEON program will be briefly described, including a multi-university workshop planned in May for students to present their OCTEON project and compete for the OCTEON Trophy.

Last Lecture

Data dependence handling

Today's Agenda

Control dependence handling

Review: How to Handle Data Dependences

- Anti and output dependences are easier to handle
 - Dependence on name, not dependence on value/dataflow
 - Get rid of them via renaming: You have seen a version of this!
- Flow dependences are more interesting
- Five fundamental ways of handling flow dependences
 - Detect and wait until value is available in register file
 - Detect and forward/bypass data to dependent instruction
 - Detect and eliminate the dependence at the software level
 - No need for the hardware to detect dependence
 - Predict the needed value(s), execute "speculatively", and verify
 - Do something else (fine-grained multithreading)
 - No need to detect

Questions to Ponder

- What is the role of the hardware vs. the software in data dependence handling?
 - Software based interlocking
 - Hardware based interlocking
 - Who inserts/manages the pipeline bubbles?
 - Who finds the independent instructions to fill "empty" pipeline slots?
 - What are the advantages/disadvantages of each?

Questions to Ponder

- What is the role of the hardware vs. the software in the order in which instructions are executed in the pipeline?
 - □ Software based instruction scheduling → static scheduling
 - □ Hardware based instruction scheduling → dynamic scheduling

More on Software vs. Hardware

- Software based scheduling of instructions → static scheduling
 - Compiler orders the instructions, hardware executes them in that order
 - Contrast this with dynamic scheduling (in which hardware will execute instructions out of the compiler-specified order)
 - How does the compiler know the latency of each instruction?
- What information does the compiler not know that makes static scheduling difficult?
 - Answer: Anything that is determined at run time
 - Variable-length operation latency, memory addr, branch direction
- How can the compiler alleviate this (i.e., estimate the unknown)?
 - Answer: Profiling

Control Dependence Handling

Control Dependence

- Question: What should the fetch PC be in the next cycle?
- Answer: The address of the next instruction
 - All instructions are control dependent on previous ones. Why?
- If the fetched instruction is a non-control-flow instruction:
 - Next Fetch PC is the address of the next-sequential instruction
 - Easy to determine if we know the size of the fetched instruction
- If the instruction that is fetched is a control-flow instruction:
 - How do we determine the next Fetch PC?
- In fact, how do we even know whether or not the fetched instruction is a control-flow instruction?

Branch Types

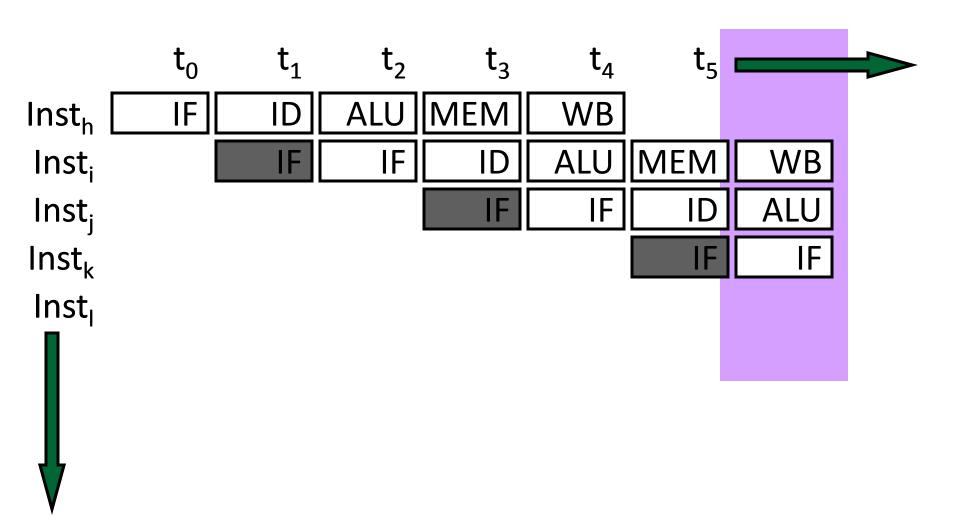
| Type | Direction at fetch time | Number of possible next fetch addresses? | When is next fetch address resolved? |
|---------------|-------------------------|--|--------------------------------------|
| Conditional | Unknown | 2 | Execution (register dependent) |
| Unconditional | Always taken | 1 | Decode (PC + offset) |
| Call | Always taken | 1 | Decode (PC + offset) |
| Return | Always taken | Many | Execution (register dependent) |
| Indirect | Always taken | Many | Execution (register dependent) |

Different branch types can be handled differently

How to Handle Control Dependences

- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
- Stall the pipeline until we know the next fetch address
- Guess the next fetch address (branch prediction)
- Employ delayed branching (branch delay slot)
- Do something else (fine-grained multithreading)
- Eliminate control-flow instructions (predicated execution)
- Fetch from both possible paths (if you know the addresses of both possible paths) (multipath execution)

Stall Fetch Until Next PC is Available: Good Idea?



This is the case with non-control-flow and unconditional br instructions!

Doing Better than Stalling Fetch ...

 Rather than waiting for true-dependence on PC to resolve, just guess nextPC = PC+4 to keep fetching every cycle

```
Is this a good guess?
What do you lose if you guessed incorrectly?
```

- ~20% of the instruction mix is control flow
 - □ ~50 % of "forward" control flow (i.e., if-then-else) is taken
 - ~90% of "backward" control flow (i.e., loop back) is taken Overall, typically ~70% taken and ~30% not taken [Lee and Smith, 1984]
- Expect "nextPC = PC+4" ~86% of the time, but what about the remaining 14%?

Guessing NextPC = PC + 4

- Always predict the next sequential instruction is the next instruction to be executed
- This is a form of next fetch address prediction and branch prediction
- How can you make this more effective?
- Idea: Maximize the chances that the next sequential instruction is the next instruction to be executed
 - Software: Lay out the control flow graph such that the "likely next instruction" is on the not-taken path of a branch
 - Hardware: ??? (how can you do this in hardware...)

Guessing NextPC = PC + 4

- How else can you make this more effective?
- Idea: Get rid of control flow instructions (or minimize their occurrence)
- How?
 - 1. Get rid of unnecessary control flow instructions → combine predicates (predicate combining)
 - 2. Convert control dependences into data dependences → predicated execution

Predicate Combining (not Predicated Execution)

- Complex predicates are converted into multiple branches
 - \Box if ((a == b) && (c < d) && (a > 5000)) { ... }
 - 3 conditional branches
- Problem: This increases the number of control dependencies
- Idea: Combine predicate operations to feed a single branch instruction instead of having one branch for each
 - Predicates stored and operated on using condition registers
 - A single branch checks the value of the combined predicate
- + Fewer branches in code → fewer mipredictions/stalls
- -- Possibly unnecessary work
 - -- If the first predicate is false, no need to compute other predicates
- Condition registers exist in IBM RS6000 and the POWER architecture

Predicated Execution

- Idea: Convert control dependence to data dependence
- Suppose we had a Conditional Move instruction...
 - □ CMOV condition, R1 ← R2
 - \square R1 = (condition == true) ? R2 : R1
 - Employed in most modern ISAs (x86, Alpha)
- Code example with branches vs. CMOVs if (a == 5) {b = 4;} else {b = 3;}

```
CMPEQ condition, a, 5;
CMOV condition, b \leftarrow 4;
CMOV !condition, b \leftarrow 3;
```

Predicated Execution

 Eliminates branches → enables straight line code (i.e., larger basic blocks in code)

Advantages

- Always-not-taken prediction works better (no branches)
- Compiler has more freedom to optimize code (no branches)
 - control flow does not hinder inst. reordering optimizations
 - code optimizations hindered only by data dependencies

Disadvantages

- Useless work: some instructions fetched/executed but discarded (especially bad for easy-to-predict branches)
- Requires additional ISA support
- Can we eliminate all branches this way?

Predicated Execution

- We will get back to this...
- Some readings (optional):
 - Allen et al., "Conversion of control dependence to data dependence," POPL 1983.
 - Kim et al., "Wish Branches: Combining Conditional Branching and Predication for Adaptive Predicated Execution," MICRO 2005.

How to Handle Control Dependences

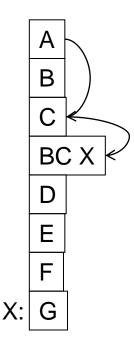
- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
- Stall the pipeline until we know the next fetch address
- Guess the next fetch address (branch prediction)
- Employ delayed branching (branch delay slot)
- Do something else (fine-grained multithreading)
- Eliminate control-flow instructions (predicated execution)
- Fetch from both possible paths (if you know the addresses of both possible paths) (multipath execution)

Delayed Branching (I)

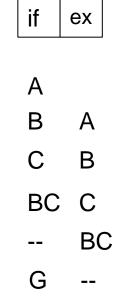
- Change the semantics of a branch instruction
 - Branch after N instructions
 - Branch after N cycles
- Idea: Delay the execution of a branch. N instructions (delay slots) that come after the branch are always executed regardless of branch direction.
- Problem: How do you find instructions to fill the delay slots?
 - Branch must be independent of delay slot instructions
- Unconditional branch: Easier to find instructions to fill the delay slot
- Conditional branch: Condition computation should not depend on instructions in delay slots → difficult to fill the delay slot

Delayed Branching (II)

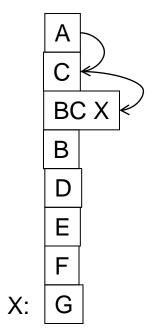
Normal code:



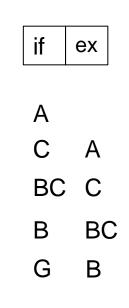
Timeline:



Delayed branch code:



Timeline:



6 cycles

5 cycles

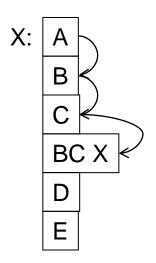
Fancy Delayed Branching (III)

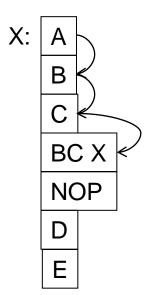
- Delayed branch with squashing
 - In SPARC
 - If the branch falls through (not taken), the delay slot instruction is not executed
 - Why could this help?

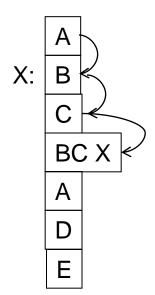
Normal code:

Delayed branch code:

Delayed branch w/ squashing:







Delayed Branching (IV)

Advantages:

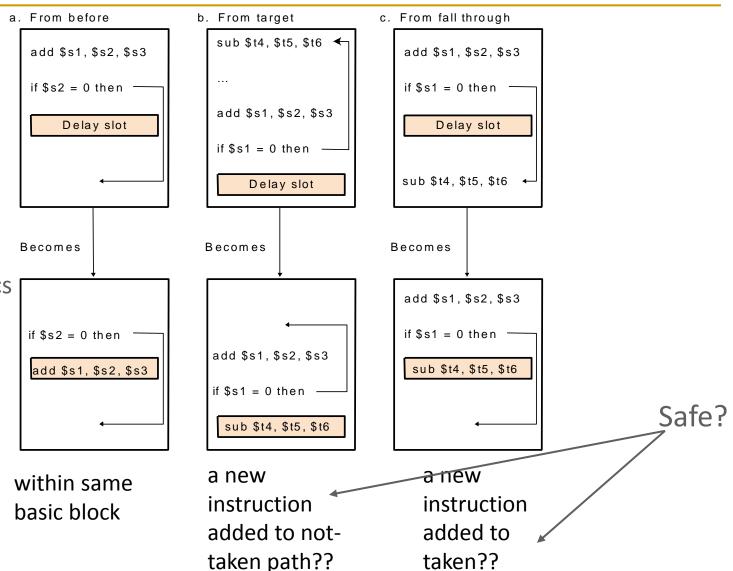
- + Keeps the pipeline full with useful instructions in a simple way assuming
- 1. Number of delay slots == number of instructions to keep the pipeline full before the branch resolves
 - 2. All delay slots can be filled with useful instructions

Disadvantages:

- -- Not easy to fill the delay slots (even with a 2-stage pipeline)
 - 1. Number of delay slots increases with pipeline depth, superscalar execution width
 - 2. Number of delay slots should be variable with variable latency operations. Why?
- -- Ties ISA semantics to hardware implementation
 - -- SPARC, MIPS, HP-PA: 1 delay slot
 - -- What if pipeline implementation changes with the next design?

An Aside: Filling the Delay Slot

reordering data
independent
(RAW, WAW,
WAR)
instructions
does not change
program semantics

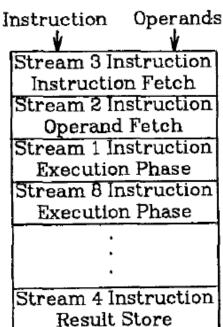


How to Handle Control Dependences

- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
- Stall the pipeline until we know the next fetch address
- Guess the next fetch address (branch prediction)
- Employ delayed branching (branch delay slot)
- Do something else (fine-grained multithreading)
- Eliminate control-flow instructions (predicated execution)
- Fetch from both possible paths (if you know the addresses of both possible paths) (multipath execution)

Fine-Grained Multithreading

- Idea: Hardware has multiple thread contexts. Each cycle, fetch engine fetches from a different thread.
 - By the time the fetched branch/instruction resolves, there is no need to fetch another instruction from the same thread
 - Branch/instruction resolution latency overlapped with execution of other threads' instructions
- + No logic needed for handling control and data dependences within a thread
- -- Single thread performance suffers
- -- Extra logic for keeping thread contexts
- Does not overlap latency if not enough threads to cover the whole pipeline



Fine-grained Multithreading

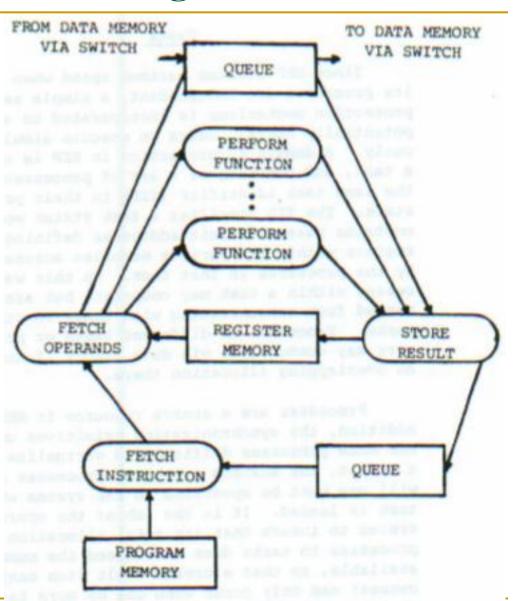
- Idea: Switch to another thread every cycle such that no two instructions from the thread are in the pipeline concurrently
- Tolerates the control and data dependency latencies by overlapping the latency with useful work from other threads
- Improves pipeline utilization by taking advantage of multiple threads
- Thornton, "Parallel Operation in the Control Data 6600," AFIPS 1964.
- Smith, "A pipelined, shared resource MIMD computer," ICPP 1978.

Fine-grained Multithreading: History

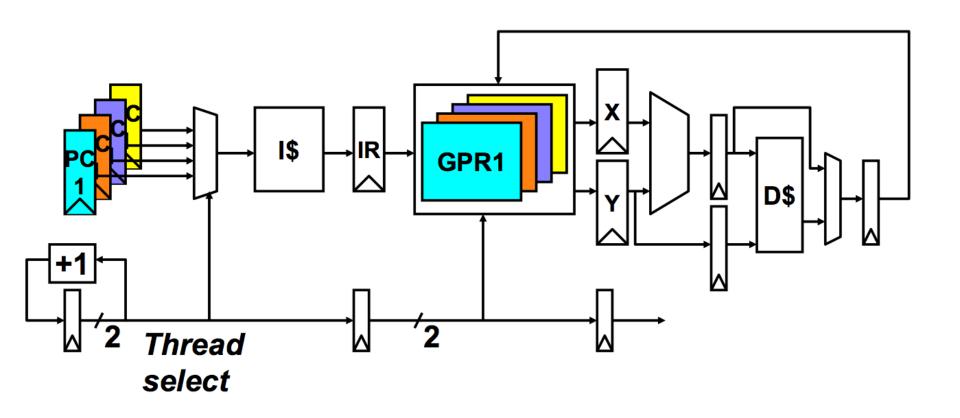
- CDC 6600's peripheral processing unit is fine-grained multithreaded
 - Thornton, "Parallel Operation in the Control Data 6600," AFIPS 1964.
 - Processor executes a different I/O thread every cycle
 - An operation from the same thread is executed every 10 cycles
- Denelcor HEP (Heterogeneous Element Processor)
 - Smith, "A pipelined, shared resource MIMD computer," ICPP 1978.
 - 120 threads/processor
 - available queue vs. unavailable (waiting) queue for threads
 - each thread can only have 1 instruction in the processor pipeline; each thread independent
 - to each thread, processor looks like a non-pipelined machine
 - system throughput vs. single thread performance tradeoff

Fine-grained Multithreading in HEP

- Cycle time: 100ns
- 8 stages → 800 ns to complete an instruction
 - assuming no memory access

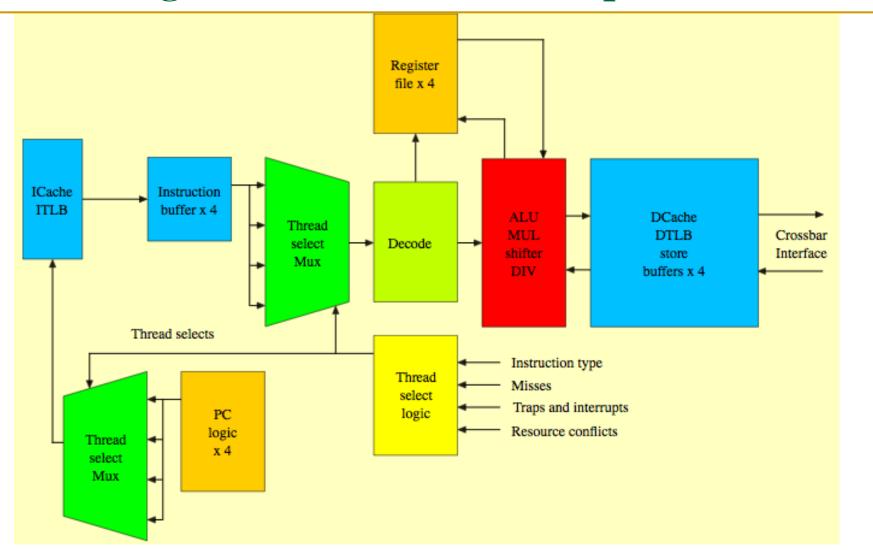


Multithreaded Pipeline Example



Slide from Joel Emer

Sun Niagara Multithreaded Pipeline



Kongetira et al., "Niagara: A 32-Way Multithreaded Sparc Processor," IEEE Micro 2005.

Fine-grained Multithreading

Advantages

- + No need for dependency checking between instructions (only one instruction in pipeline from a single thread)
- + No need for branch prediction logic
- + Otherwise-bubble cycles used for executing useful instructions from different threads
- + Improved system throughput, latency tolerance, utilization

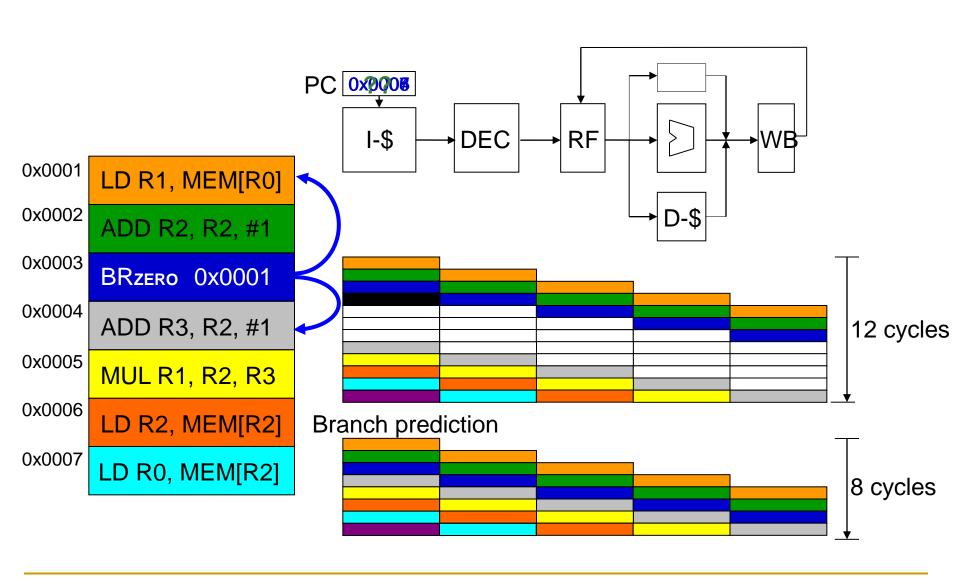
Disadvantages

- Extra hardware complexity: multiple hardware contexts, thread selection logic
- Reduced single thread performance (one instruction fetched every N cycles)
- Resource contention between threads in caches and memory
- Some dependency checking logic between threads remains (load/store)

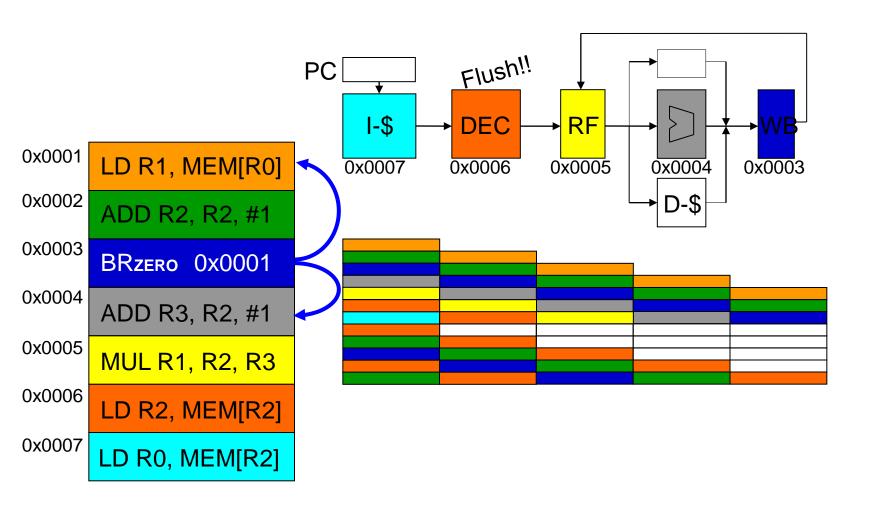
How to Handle Control Dependences

- Critical to keep the pipeline full with correct sequence of dynamic instructions.
- Potential solutions if the instruction is a control-flow instruction:
- Stall the pipeline until we know the next fetch address
- Guess the next fetch address (branch prediction)
- Employ delayed branching (branch delay slot)
- Do something else (fine-grained multithreading)
- Eliminate control-flow instructions (predicated execution)
- Fetch from both possible paths (if you know the addresses of both possible paths) (multipath execution)

Branch Prediction: Guess the Next Instruction to Fetch

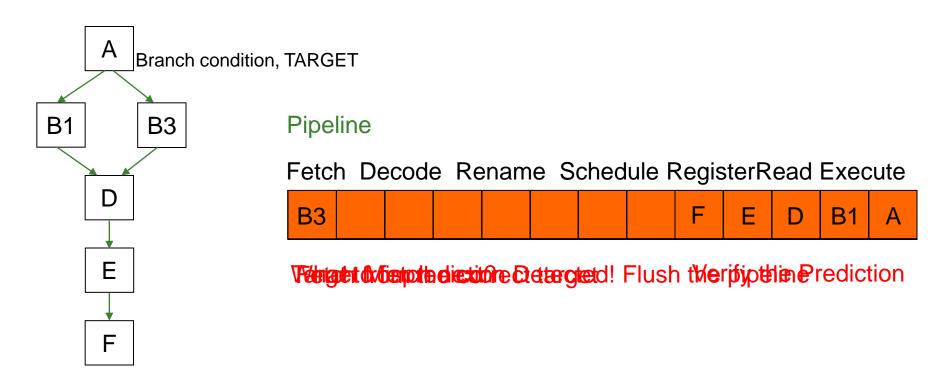


Misprediction Penalty



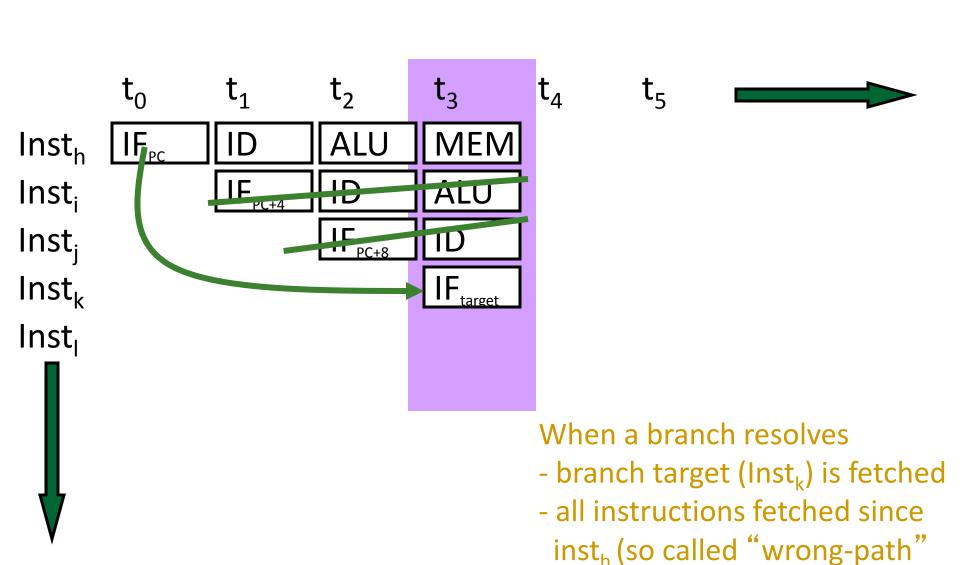
Branch Prediction

- Processors are pipelined to increase concurrency
- How do we keep the pipeline full in the presence of branches?
 - Guess the next instruction when a branch is fetched
 - Requires guessing the direction and target of a branch



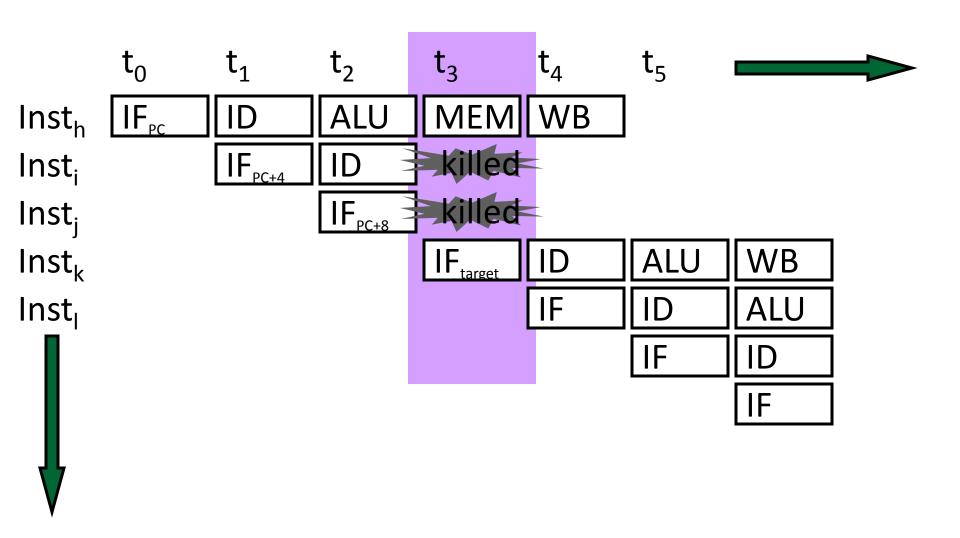
Branch Prediction: Always PC+4

Inst_h is a branch



instructions) must be flushed

Pipeline Flush on a Misprediction



Inst_h is a branch

Performance Analysis

correct guess ⇒ no penalty

~86% of the time

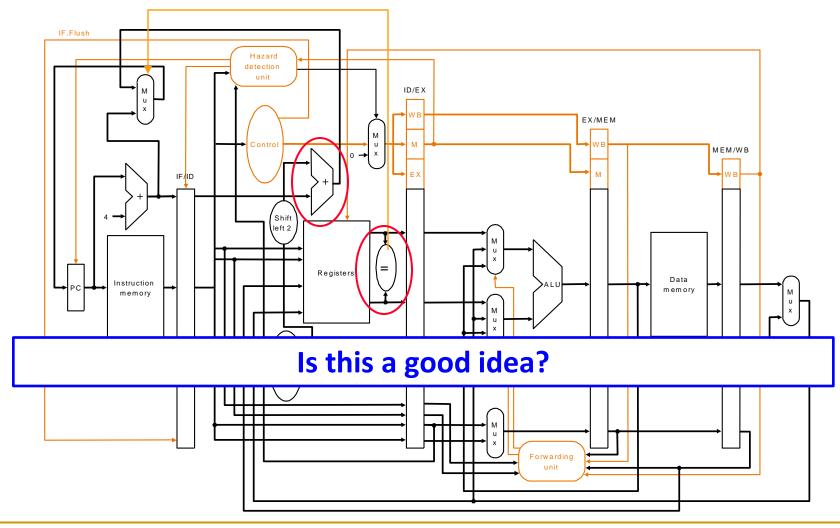
- incorrect guess \Rightarrow 2 bubbles
- Assume
 - no data hazards
 - 20% control flow instructions
 - □ 70% of control flow instructions are taken

probability of penalty for a wrong guess a wrong guess

Can we reduce either of the two penalty terms?

Reducing Branch Misprediction Penalty

Resolve branch condition and target address early

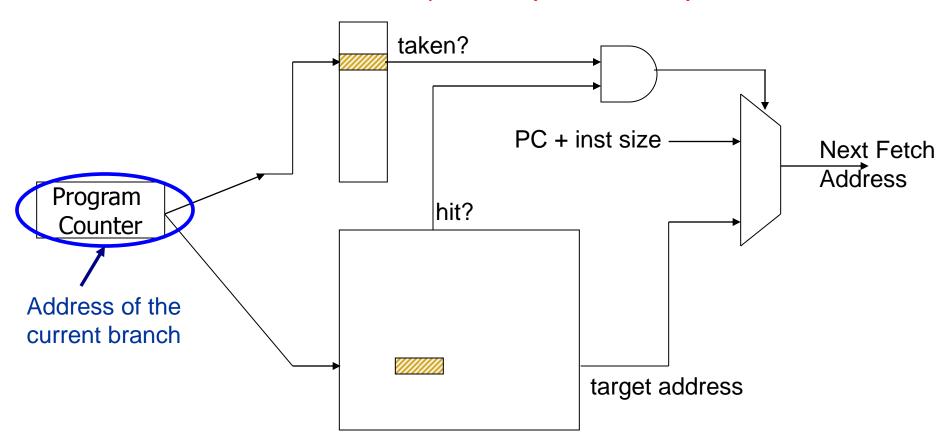


Branch Prediction (Enhanced)

- Idea: Predict the next fetch address (to be used in the next cycle)
- Requires three things to be predicted at fetch stage:
 - Whether the fetched instruction is a branch
 - (Conditional) branch direction
 - Branch target address (if taken)
- Observation: Target address remains the same for a conditional direct branch across dynamic instances
 - Idea: Store the target address from previous instance and access it with the PC
 - Called Branch Target Buffer (BTB) or Branch Target Address
 Cache

Fetch Stage with BTB and Direction Prediction

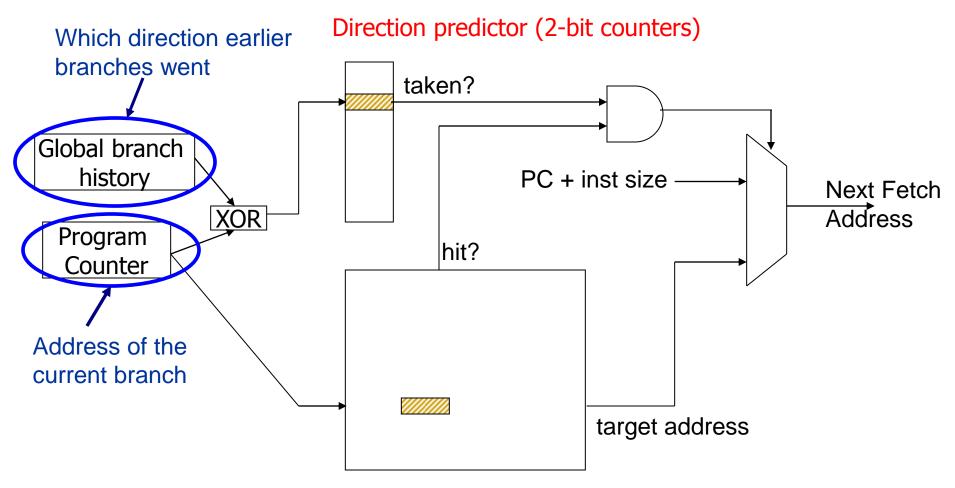
Direction predictor (2-bit counters)



Cache of Target Addresses (BTB: Branch Target Buffer)

Always taken CPI = [1 + (0.20*0.3)*2] = 1.12 (70% of branches taken)

More Sophisticated Branch Direction Prediction



Cache of Target Addresses (BTB: Branch Target Buffer)

Simple Branch Direction Prediction Schemes

- Compile time (static)
 - Always not taken
 - Always taken
 - BTFN (Backward taken, forward not taken)
 - Profile based (likely direction)
- Run time (dynamic)
 - Last time prediction (single-bit)

More Sophisticated Direction Prediction

- Compile time (static)
 - Always not taken
 - Always taken
 - BTFN (Backward taken, forward not taken)
 - Profile based (likely direction)
 - Program analysis based (likely direction)
- Run time (dynamic)
 - Last time prediction (single-bit)
 - Two-bit counter based prediction
 - Two-level prediction (global vs. local)
 - Hybrid